Tight Lower Bounds for Weighted Matroid Problems

Anonymous

July 9, 2023

Abstract

1 Introduction

Matroids are simple combinatorial structures which provide a unified abstractions for independence systems such as linear independence in a vector space and cycle-free subsets of edges (of a given graphs). While serving as a generic abstraction, matroids poses sufficient combinatorial properties to allow the development of efficient algorithms. Such algorithm include the classic greedy which finds a maximum weight independent set of a matroid, Lawler's [13] algorithm for matroid intersection and the more modern pipage rounding technique [2]. As such, a significant attention has been given for various forms of combinatorial matroid optimization and decision problems.

In this paper we give hardness results for several fundamental weighted matroid problems. We show that there is no *fully polynomial time approximation scheme* (FPTAS) for the *budgeted matroid independent set problem*, indicating the *efficient polynomial time approximation scheme* (EPTAS) of Doron-Arad, Kulik and Shachnai [9] is the best possible result for the problem. Additionally, we show that there is no pseudo-polynomial algorithm for *exact weight matroid independent set* (EMI), indicating the algorithm of Camerini, Galbiati and Maffioli [3] for *repersented* matroids cannot be generalized to arbitrary matroids. Similarly, we show there is no Fully PTAS for *constrained minimum basis of a matroid* and *knapsack cover with a matroid*, indicating the existing Efficient PTASs are optimal.

For all the above problems we provide unconditional lower bounds in the *oracle* model where the independent sets of the matorid can only be accessed through a membership oracle. We complement these results with lower bounds based on standard complexity assumptions which provide the same results even if the matroid is encoded as part of the instance. All our lower bounds are based on a specifically structured family of *paving matroids*.

1.1 Basic Definitions

Our results hold also for both randomized and deterministic algorithms. In this context, a *randomized* algorithm $\mathcal{A}$ for a decision problem $\mathcal{D}$, is an algorithm that given an instance I of $\mathcal{D}$ returns the following.

- If I is a no-instance, the algorithm returns it is a no-instance with probability (w.p.) 1.
- If I is a yes-instance, the algorithm returns that I is a yes-instance w.p. at least $\frac{1}{2}$.

We also consider optimization problems that do not necessarily have a solution; however, in all of these problems deciding if a solution exists is polynomially solvable. For an instance I of an optimization problem P, let $\mathrm{OPT}(I)$ be the optimum value of a solution for I. For some $\rho \geq 1$, a randomized ρ-approximation algorithm $\mathcal{A}$ for an optimization problem P, is an algorithm that given an instance I of P returns the following.

- If I does not have a solution, the algorithm returns that I does not have a solution w.p. 1.
- If I has a solution, the algorithm returns w.p. at least $\frac{1}{2}$ a solution S for I of value v satisfying:
  - $v \geq \frac{\mathrm{OPT}(I)}{\rho}$ if P is a maximization problem.
  - $v \leq \rho \cdot \mathrm{OPT}(I)$ if P is a minimization problem.

For $\rho = 1$, we say that $\mathcal{A}$ is a randomized exact algorithm for P. For both decision and optimization problems, if the second condition, respectively holds w.p. 1, we say that algorithm $\mathcal{A}$ is deterministic.

Our main results indicate what is the best possible type of an approximation scheme a problem admits. For an instance I of a combinatorial optimization problem P, let $|I|$ be the encoding size

of I. A (randomized) *polynomial-time approximation scheme* (PTAS) for P is a family of algorithms $(A_\varepsilon)_{\varepsilon>0}$ such that, for any $\varepsilon > 0$, A_ε is a polynomial-time (randomized) $(1+\varepsilon)$-approximation algorithm for P. An *Efficient PTAS* (EPTAS) is a PTAS $(A_\varepsilon)_{\varepsilon>0}$ with running time of the form $f\left(\frac{1}{\varepsilon}\right) \cdot |I|^{O(1)}$, where f is an arbitrary computable function; In addition, $(A_\varepsilon)_{\varepsilon>0}$ is a (randomized) *Fully PTAS* (FPTAS) if the running time of A_ε is of the form $\left(\frac{|I|}{\varepsilon}\right)^{O(1)}$.

All the problems considered in this paper involve *matroids*.

Definition 1.1. *A* matroid *is a set system $(E, \mathcal{I})$, where E is a finite set and $\mathcal{I} \subseteq 2^E, \mathcal{I} \neq \emptyset$ such that*

- *The* hereditary property*: for all $A \in \mathcal{I}$ and $B \subseteq A$ it holds that $B \in \mathcal{I}$.*
- *The* exchange property*: for all $A, B \in \mathcal{I}$ where $|A| > |B|$ there is $e \in A \setminus B$ such that $B + e \in \mathcal{I}$.*

An independent set $S \in \mathcal{I}$ is a basis *if $|S| = \max_{S' \in \mathcal{I}} |S'|$.*

Several books and book chapters have been dedicated for matroids from both the combinatorial and algorithmic point of view. We refer the reader to [16] and [18] for additional information about matroids.

In the common computational setting, the set system $(E, \mathcal{I})$ is not given explicitly as part of the input due to the large encoding size required for $\mathcal{I}$. The standard computational model that bypasses the issue is the *membership oracle* model; in this model, testing whether a given subset $S \subseteq E$ belongs to $\mathcal{I}$ is determined in a single query to the membership oracle. The oracle model enables us to obtain unconditional information-theoretic based lower bounds for the running time of algorithms for matroid related problems. While these lower bounds hold in the oracle model, one may suspect that the lower bounds do not hold for any version of the matroid problem in which the matroid is encoded as part of the input. To rule this out, we also consider an alternative model in which the matroid $(E, \mathcal{I})$ is encoded as part of the input and membership can be decided in polynomial time. We further elaborate on this model in Section 5.

1.2 Our Results

In this paper we give several hardness results for weighted matroid problem. A summary of our results is given in Table 1.

Problem	Previous Results	This paper
Budgeted Matroid Independent Set	Efficient PTAS [9]	No Fully PTAS
Budgeted Matroid Intersection	Efficient PTAS [9]	No Fully PTAS
Exact Weight Matroid Independent Set	A pseudo polynomial time algorithm for *representable* matroids [3]	No pseudo polynomial algorithm for general matroids.
Constraint Minimum Basis of a Matroid	Efficient PTAS [12]	No Fully PTAS
Knapsack Cover with a Matroid	Efficient PTAS [12] (implicit)	No Fully PTAS

Table 1: A summary of our results.

1.2.1 The Budgeted Matroid Independent Set

Our first result rules out the existence of a Fully PTAS for *budgeted matroid independent set*.

Budgeted Matroid Independent Set (BM)

Instance	$I = (E, \mathcal{I}, c, p, B)$, where $(E, \mathcal{I})$ is a matroid, $c : E \to \mathbb{N}$ is the cost function, $p : E \to \mathbb{N}$ is the profit function and $B \in \mathbb{N}$ is the budget
Solution	$S \in \mathcal{I}$ such that $c(S) \leq B$.
Objective	Find a solution S with maximum profit $p(S)$.
Model	The collection of independent sets $\mathcal{I}$ is only accessible via a membership oracle; the encoding cost of the instance I is the encoding cost of (E, c, p, B), the instance without $\mathcal{I}$.

The problem generalizes the 0/1-knapsack problem for which a Fully PTAS is known since the 1970's [14]. A Fully PTAS is also known for other special cases of BM such as *multiple choice knpasack* [14] and 1.5-*dimensional knapsack* [4].

A PTAS for BM was first given in [1] as a special case of *budgeted matroid intersection*, a generalization of BM with two matroids. A PTAS for the problem can also be derived as a special case of the results in [11, 6] which provide PTASs for a multi-budgeted variants of BM. An *Efficient* PTAS for BM was recently given in [9], leaving the existence of a *Fully* PTAS for the problem as an open question. We give an answer to this question by showing that there is no Fully PTAS for the problem.

Theorem 1.2. *There is no randomized fully polynomial time approximation scheme (FPTAS) for budgeted matroid independent set (BM).*

Observe Theorem 1.2 does not use any complexity related assumptions. The proof is information theoretic and exploits the limitations of the oracle model. As already mentioned, we also provide a variant of Theorem 1.2 which rules out the existence of a Fully PTAS for a version of BM in which the matroid belongs to a specific family of matroids and is encoded as part of the input for the algorithm.

Theorem 1.2 also rules out the existence of a Fully PTAS for the following generalization of BM.

Budgeted Matroid Intersection (BMI)

Instance	$I = (E, \mathcal{I}_1, \mathcal{I}_2, c, p, B)$, where $(E, \mathcal{I}_1)$ and $(E, \mathcal{I}_2)$ are matroids, $c : E \to \mathbb{N}$ is the cost function, $p : E \to \mathbb{N}$ is the profit function and $B \in \mathbb{N}$ is the budget.
Solution	$S \in \mathcal{I}_1 \cap \mathcal{I}_2$ such that $c(S) \leq B$.
Objective	Find a solution S of I with maximum profit $p(S)$.
Model	Both $\mathcal{I}_1$ and $\mathcal{I}_2$ are only accessible via a membership oracle; the encoding cost of the instance I is the encoding cost of (E, c, p, B), the instance without $\mathcal{I}_1$ and $\mathcal{I}_2$.

That is, BMI is a variant of BM in which we are given *two* matroids $(E, \mathcal{I}_1)$ and $(E, \mathcal{I}_2)$, and a solution has to be an independent set of *both* matroids. A PTAS for BMI was first given by [1]. The result was recently improved to an Efficient PTAS [8]. Since BMI is a generalization of BM (e.g., take $\mathcal{I}_1 = \mathcal{I}_2$), Theorem 1.2 implies the following.

Corollary 1.3. *There is no fully polynomial time approximation scheme (FPTAS) for budgeted matroid intersection (BMI).*

That is, the Efficient PTAS of [8] for BMI is the best possible result for the problem.

1.2.2 Exact Weight Matroid Independent Set

Our next result shows that there is no pseudo-polynomial time algorithm for exact weight matroid independent set.

Exact Matroid Independent Set (EMI)

Instance	$I = (E, \mathcal{I}, w, k, T)$, where $(E, \mathcal{I})$ is a matroid along with a membership oracle, $w : E \to \mathbb{N}$ gives the elements weights, $k \in \mathbb{N}$ is the target cardinality, and $T \in \mathbb{N}$ is the target weight.
Objective	Decide if there is $S \in \mathcal{I}$ such that $\lvert S\rvert = k$ and $w(S) = T$.
Model	The independent sets $\mathcal{I}$ are only accessible via a membership oracle.

We observe that EMI is a generalization of k-Subset-Sum (the special case of EMI in which $\mathcal{I} = 2^E$), and hence we do expect the problem to be solvable in polynomial time. Instead, we are looking for pseudo-polynomial algorithm whose running time has polynomial dependence in the encoding size and in T. Indeed, the special case of EMI in which the matroid is *representable* admit such an algorithm (see details below). Our second result rules our the existence of a psuedo-polynomial algorithm for EMI, even if we additionally allow the running time to be parameterized in k.

Theorem 1.4. *There is no randomized algorithm for EMI which runs in time $f(k) \cdot (T \cdot |I|)^{O(1)}$ for any function f, where the instance is $I = (E, \mathcal{I}, w, k, T)$ and $|I|$ is the encoding cost of (E, w, k, T).*

As in Theorem 1.2, Theorem 1.4 does not require any complexity related assumptions. As a special case, Theorem 1.4 implies there is no algorithm for EMI which runs in time $(T \cdot |I|)^{O(1)}$.

A *representable matroid* over a field $\mathbb{F}$ is a matroid $(E, \mathcal{I})$ where $E \subseteq \mathbb{F}^\ell$ for some $\ell \in \mathbb{N}$ and

$$\mathcal{I} = \{S \subseteq E \mid S \text{ is linearly independent in } \mathcal{F}^\ell\}.$$

While it is easy to show the above construction always leads to a matroid, not all matroids are representable (see, e.g., [16]) (Ariel :verify this claim indeed appears there). Assuming the elements of $\mathbb{F}$ can be encoded, representable matroids can be easily encoded as part a problem instance by simply encoding the vectors in E as part of the input.

In *representable* EMI the input is a representable matroid $(E, \mathcal{I})$ over the rationals, weight function $w : E \to \mathbb{N}$, and target weight $T \in \mathbb{N}$. The objective is to determine if there exists an independent set of the matroid of weight exactly T, that is, $S \in \mathcal{I}$ such that $w(S) = T$. The result in [3] together with [15] provide a randomized algorithm for representable EMI which runs in polynomial time in the encoding of the instance and in T. We note the result in [3] requires the solution to be a basis and is given for the more generic *exact parity basis* problem. The requirement that the solution is a basis can be easily circumvented using the truncation technique of [15] for representable matroids. Theorem 1.4 implies this result cannot be generalized to arbitrary matroids.

Theorem 1.4 can be used to show that an algorithm of Fomin et al. [10] which solves a matroid problem for representable matroids cannot be generalized to arbitrary matroids. We give the details in Section 7. This additional application uses the fact that Theorem 1.4 also rules out algorithms with parameterized running time in k.

1.2.3 Constrained Minimum Basis of a Matroid

Our next result shows there is no Fully PTAS for a minimization variant of BM in which the solution must be a basis of the matroid. Recall that a basis of a matroid $(E, \mathcal{I})$ is an independent set $S \in \mathcal{I}$ such that $|S| = \max_{S' \in \mathcal{I}} |S'|$.

Constrained Minimum Basis of a Matroid (CMB)

Instance	$I = (E, \mathcal{I}, c, w, B)$ where $(E, \mathcal{I})$ is a matroid, $c : E \to \mathbb{N}$ is the cost function, $w : E \to \mathbb{N}$ is the weight function, and $B \in \mathbb{N}$ is the budget.
Solution	A **basis** S of the matroid $(E, \mathcal{I})$ such that $c(S) \leq B$.
Objective	Find a solution of minimum weight $w(S)$.
Model	The independent sets $\mathcal{I}$ are only accessibly via a membership oracle. The encoding cost of the instance is the encoding cost of (E, ℓ, c, B), the instance without $\mathcal{I}$.

Observe that there are CMB instances for which there is no solution. This occurs, for example, if $\mathcal{I} = \{S \subseteq E \mid |S| \leq k\}$, $B = k - 1$ and $c(e) = 1$ for all $e \in E$. However, the minimum cost basis of $(E, \mathcal{I})$ can be computed in polynomial time, hence , it can be determined in polynomial time if the solution set of an instance is empty.

A *graphical* matroid $(E, \mathcal{I})$ is a matroid for which there is a graph $G = (V, E)$ and it holds that

$$\mathcal{I} = \{S \subseteq E \mid S \text{ is a forest}\}.$$

That is, the independent sets of the matroid are cycle-free subsets of edges of the graph G. A PTAS for the special case of CMB in which the matroid is a graphical matroid was given by Ravi and Goemans [17]. This results was improved to an Efficient PTAS by Hassin and Levin [12]. As indicated in [12], the algorithm [12] is actually an Efficient PTAS for the general CMB problem. The next theorem indicates this is the best possible result for the problem.

Theorem 1.5. *There is no fully polynomial time approximation scheme (FTPAS) for constraint minimum basis of a matroid (CMB).*

1.2.4 Knapsack Cover with a Matroid

Our final result rules our the existence of a Fully PTAS for a *coverage* variant of knapsack.

Knapsack Cover with a Matroid (KCM)

Instance	$I = (E, \mathcal{I}, c, d, D)$, where $(E, \mathcal{I})$ is a matroid, $c : E \to \mathbb{N}$ is the cost function costs, $d : E \to \mathbb{N}$ is the size function, and $D \in \mathbb{N}$ is the demand.
Solution	A set $S \in \mathcal{I}$ such that $d(S) \geq D$.
Objective	Find a solution S of minmum cost $c(S)$.
Model	The independent sets $\mathcal{I}$ are only accessible via a membership oracle. The encoding cost of the instance is the encoding cost of (E, c, d, D), the instance without $\mathcal{I}$.

In [5] Chakaravarthy et al. provided a PTAS for KCM. It seems, however, that an Efficient PTAS for KCM can be attained by adapting the approach of Hassin and Levin [12] for constraint minimum spanning trees to the setting KCM (we omit the details). We show that an Efficient PTAS is indeed the best possible approximation scheme for KCM.

Theorem 1.6. *There is no random fully polynomial time approximation scheme (FPTAS) for knapsack cover with a matroid (KCM).*

1.2.5 Hardness Results for Encoded Matroids

The lower bounds in Theorems 1.2 and 1.4 to 1.6 all rely on the oracle model. In Section 6 we show variants of the Theorems which hold in a setting where the matroid is encoded as part of the input.

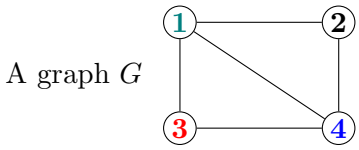

$$\mathcal{J}_{4,2} = \{\emptyset, \{\mathbf{1}\}, \{\mathbf{2}\}, \{\mathbf{3}\}, \{\mathbf{4}\}\}, \quad \mathcal{K}_{4,2,5} = \{\{\mathbf{2},\mathbf{4}\}, \{\mathbf{2},\mathbf{1}\}, \{\mathbf{4},\mathbf{3}\}, \{\mathbf{3},\mathbf{1}\}\}, \quad \mathcal{L}_{4,2,5}(\Pi) = \{\{\mathbf{2},\mathbf{3}\}\}$$

Figure 1: The independent sets of the Π-matroid $M_{n,k,\alpha}(\Pi)$, with parameters $n = 4$, $k = 2$, $\alpha = 5$. The secret set Π contains all independent sets in the graph G, where $\{2, 3\}$ is the only independent set in G with k elements.

These theorems do assume P $\neq$ NP. A proper statement of these results require several technical definitions about encoding of matroids. We therefore defer the statement of these theorems to Section 6.

1.3 Technical Overview

2 Preliminaries

2.1 Notations

For the simplicity of the notation, for any set A and some element a, let $A + a$, $A - a$ be $A \cup \{a\}$ and $A \setminus \{a\}$, respectively. In addition, for all $k \in \mathbb{N}$ let $[k]$ be the set $\{1, 2, \ldots, k\}$. For any function $f : A \to \mathbb{R}$ and $C \subseteq A$ let $f(C) = \sum_{\ell \in C} f(\ell)$. Finally, given the set of numbers $[n]$ for some $n \in \mathbb{N}$ and $S \subseteq [n]$, we use $\mathsf{sum}(S) = \sum_{i \in S} i$. We use $\mathrm{OPT}(I)$ to denote the optimum value of an instance I of some optimization problem.

2.2 Randomized Algorithms and Approximation Schemes

2.3 Matroids and Membership Oracles

3 The Π-Matroid Family

In this section, we formally define the Π-matroid problem and prove an unconditional lower bound for solving it in the oracle model. In simple words, the matroids in this section are a family of matroids defined over the following parameters. The first parameter $n \in \mathbb{N}$ is the number of elements and the ground set is simply $[n]$. The second parameter, $k \in [n]$, is the rank of the matroid.[1] Furthermore, α is a target value, which is usually equals to $\mathsf{sum}(S)$ for some $S \subseteq [n]$. Finally, the last parameter is a family of subsets $\Pi \subseteq 2^{[n]}$, where $2^{[n]} = 2^{\{1,\ldots,n\}}$.

The set Π is called the *secret set* because finding Π is possible only via recurrent queries to the membership oracle of the matroid. Since Π can have an arbitrary structure, this may require a brute force enumeration. Generally, our use of the Π-matroid family is to reduce the hardness of finding $S \in \Pi$ to the problem of finding an independent set of the matroid with unique properties w.r.t. the parameters k, α, and Π.

Definition 3.1. *Let $n, k, \alpha \in \mathbb{N}_{>0}$. For some $\Pi \subseteq 2^{[n]}$, define the* **Π-matroid** *on n, k, α as $M_{n,k,\alpha}(\Pi) = ([n], \mathcal{I}_{n,k,\alpha}(\Pi))$, where*

$$\mathcal{I}_{n,k,\alpha}(\Pi) = \mathcal{J}_{n,k} \cup \mathcal{K}_{n,k,\alpha} \cup \mathcal{L}_{n,k,\alpha}(\Pi)$$

[1] The rank of the matroid is the maximum cardinality of an independent set of the matroid.

and $\mathcal{J}_{n,k}, \mathcal{K}_{n,k,\alpha}, \mathcal{L}_{n,k,\alpha}(\Pi)$ *are defined as follows.*

$$\begin{aligned} \mathcal{J}_{n,k} &= \{S \subseteq [n] \mid |S| < k\} \\ \mathcal{K}_{n,k,\alpha} &= \{S \subseteq [n] \mid |S| = k, \mathsf{sum}(S) \neq \alpha\} \\ \mathcal{L}_{n,k,\alpha}(\Pi) &= \{S \subseteq [n] \mid |S| = k, \mathsf{sum}(S) = \alpha, S \in \Pi\}. \end{aligned} \tag{1}$$

In simple words, $\mathcal{J}_{n,k}$ contains all subsets with strictly less than k elements; $\mathcal{K}_{n,k,\alpha}$ contains all subsets of cardinality k whose total sum is not α; finally, $\mathcal{L}_{n,k,\alpha}(\Pi)$ contains all subsets of cardinality k and total sum α which also belong to Π. Observe that $\mathcal{K}_{n,k,\alpha} \cup \mathcal{L}_{n,k,\alpha}(\Pi)$ is the set of bases of the matroid. See Figure 1 for an example of a member of the Π-matroid family. When clear from the context, we omit some/all of the parameters n, k, α. We first convince the reader that this is a family of matroids.

Lemma 3.2. *For every* $n, k, \alpha \in \mathbb{N}_{>0}$*, and* $\Pi \subseteq 2^{[n]}$ *it holds that* $M(\Pi)$ *is a matroid.*

Proof. First, note that $\emptyset \in \mathcal{J}$ since $0 < k$; therefore, $\emptyset \in \mathcal{I}(\Pi)$. For the hereditary property, let $A \in \mathcal{I}(\Pi)$. For all $B \subset A$ it holds that $|B| < k$; thus, $B \in \mathcal{J}$ and it follows that $B \in \mathcal{I}(\Pi)$. For the exchange property, let $A, B \in \mathcal{I}(\Pi)$ such that $|A| > |B|$. We consider the following cases.

1. $|B| < k - 1$. Then, for all $e \in A \setminus B$ it holds that $|B + e| < k - 1 + 1 = k$; thus, $B + e \in \mathcal{J}$ and it follows that $B + e \in \mathcal{I}(\Pi)$. Note that there is such $e \in A \setminus B$ because $|A| > |B|$.

2. $|B| = k - 1$ and $|A| = k$. We consider two sub cases.

   (a) $B \subseteq A$. Then, because $|A| > |B|$ there is $e \in A \setminus B$; thus, $B + e = A$ (because $|B| = k-1$ and $|A| = k$). Hence, since $A \in \mathcal{I}(\Pi)$ it follows that $B + e \in \mathcal{I}(\Pi)$.

   (b) $B \not\subseteq A$. Then, since $|B| = k - 1$ and $|A| = k$ it follows that $|B \cap A| < |B| = k - 1$; thus, $|A \setminus B| = |A| - |A \cap B| > k - (k - 1) = 1$. Therefore, there are $e, f \in A \setminus B$ such that $e \neq f$. It follows that there is $g \in \{e, f\}$ such that $\mathsf{sum}(B + g) = \mathsf{sum}(B) + g \neq \alpha$. We conclude by (1) that $B + g \in \mathcal{K}$ and it follows that $B + g \in \mathcal{I}(\Pi)$ as required.

■

Using the Π-matroid family we define the following Π-*matroid problem*, where the goal is to decide if $\mathcal{L}(\Pi) \neq \emptyset$, where Π is some unknown fixed family of subsets of the ground set.

Definition 3.3. *The* **Π-matroid problem***: Given* $n, k \in \mathbb{N}_{>0}$*,* $\alpha \in [n^2]$*, and a membership oracle to* $M_{n,k,\alpha}(\Pi)$*, for some* $\Pi \subseteq 2^{[n]}$*, decide if* $\mathcal{L}_{n,k,\alpha}(\Pi) \neq \emptyset$*.*

Observe that for some $n, k, \alpha \in \mathbb{N}_{>0}$, $\Pi \subseteq 2^{[n]}$, and $S \subseteq [n]$, if $S \in \mathcal{L}_{n,k,\alpha}(\Pi)$ then $S \in \Pi$ as well. Thus, by (1), finding an independent set S in the matroid with $|S| = k$ and $\mathsf{sum}(S) = \alpha$ is at least as hard as finding $S \in \Pi$. This gives the intuition for the hardness of the Π-matroid problem.

We use $I = (n, k, \alpha, M_{n,k,\alpha}(\Pi))$ to denote an instance of Π-matroid; we refer to n as the *size* of the instance I and to k as the *parameter* of the instance. A *solution* of I is a set $S \in \mathcal{L}(\Pi)$. It is important to note that the input in the Π-matroid problem does not contain a clear description of Π. Instead, finding Π requires a sequence of queries to the membership oracle of $M_{n,k,\alpha}(\Pi)$. The following hardness result relies on taking a family Π that consists of a single set. Roughly speaking, to decide such a Π-matroid instance $I = (n, k, \alpha, M_{n,k,\alpha}(\Pi))$, an algorithm for the problem must iterate over all (exponentially many) subsets $S \subseteq [n]$ such that $|S| = k$ and $\mathsf{sum}(S) = \alpha$.

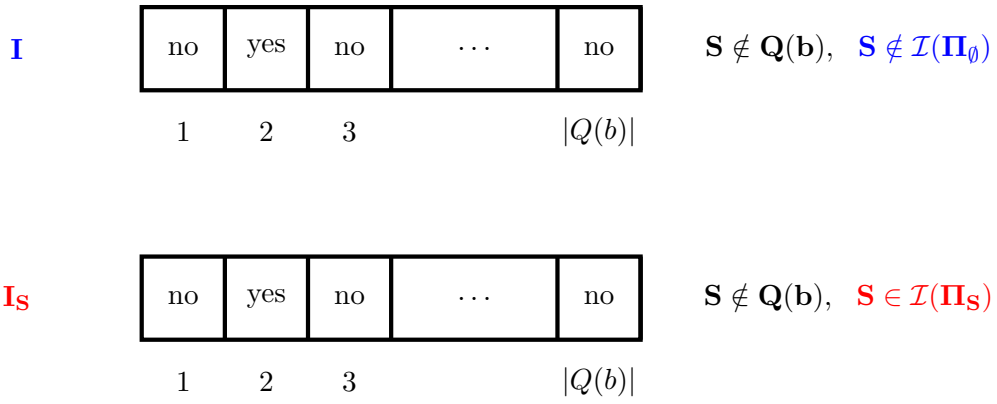

Figure 2: An illustration of the proof of Lemma 3.4. The figure presents the sequences of queries to the membership oracles by the algorithm on the instances I and I_S for some string of bits b such that $S \notin Q(b)$. The label "yes" ("no") means that the queried set is (not) independent in the matroid. The only query that distinguishes between I and I_S is on the set S, which is not queried; thus, the algorithm returns the same output for I and I_S.

Lemma 3.4. *Every randomized algorithm that decides the* Π*-matroid problem runs in time* $\Omega\left(\frac{\binom{n}{k}}{n^3}\right)$*, where* n *is the size of the instance and* k *is the parameter of the instance.*

Proof. Assume towards a contradiction that there is a randomized algorithm $\mathcal{A}$ that decides the Π-matroid problem in time $\Omega\left(\frac{\binom{n}{k}}{n^3}\right)$, where n is the size of the instance and k is the parameter of the instance. Thus, there is a constant $c \in \mathbb{N}$ such that for every $n, k \in \mathbb{N}_{>0}$ and for every Π-matroid instance I of size n and parameter k, algorithm $\mathcal{A}$ decides I in time bounded by $c \cdot \frac{\binom{n}{k}}{n^3}$. Let $\tilde{n} \in \mathbb{N}_{>0}$ such that $\tilde{n}^3 > 2c \cdot \tilde{n}^2$ and let $\tilde{k} \in \mathbb{N}_{>0}$. For simplicity, let $t = c \cdot \frac{\binom{\tilde{n}}{\tilde{k}}}{\tilde{n}^3}$. For some $\alpha \in \mathbb{N}_{>0}$, consider the following set of all subsets of $[\tilde{n}]$ with cardinality $\tilde{k}$ and sum exactly α:

$$\mathcal{F}_{\tilde{n},\tilde{k},\alpha} = \left\{S \subseteq [\tilde{n}] \;\middle|\; |S| = \tilde{k}, \mathsf{sum}(S) = \alpha\right\}.$$

Claim 3.5. *There is* $\alpha \in \left[\tilde{n}^2\right]$ *such that* $|\mathcal{F}_{\tilde{n},\tilde{k},\alpha}| \geq \frac{\binom{\tilde{n}}{\tilde{k}}}{\tilde{n}^2}$.

Proof. Observe that each $S \in 2^{[\tilde{n}]}$ where $|S| = \tilde{k}$ satisfies that

$$1 \leq \mathsf{sum}(S) \leq |S| \cdot \max_{i \in [\tilde{n}]} i = \tilde{n}^2.$$

Thus, there are at most $\tilde{n}^2$ possibilities for choosing the value α, where $\alpha = \mathsf{sum}(S)$ for some $S \in 2^{[\tilde{n}]}, |S| = \tilde{k}$. Moreover, there are $\binom{\tilde{n}}{\tilde{k}}$ subsets of $[\tilde{n}]$ with cardinality $\tilde{k}$. Therefore, by the pigeonhole principle there is $\alpha \in \left[\tilde{n}^2\right]$ such that $|\mathcal{F}_{\tilde{n},\tilde{k},\alpha}| \geq \frac{\binom{\tilde{n}}{\tilde{k}}}{\tilde{n}^2}$. □

Formally, given a Π-matroid instance $I = (\tilde{n}, \tilde{k}, \alpha, M_{\tilde{n},\tilde{k},\alpha}(\Pi))$, the randomized algorithm $\mathcal{A}$ produces a random string of bits $\bar{b} \in \{0,1\}^t$ and performs a sequence of queries to the membership oracle of $M_{\tilde{n},\tilde{k},\alpha}$ based on $\bar{b}, \tilde{n}, \tilde{k}, \alpha$, and the results of the previous queries; then, the algorithm decides the given instance based on the queries. Let $\alpha \in \left[\tilde{n}^2\right]$ such that $|\mathcal{F}_{\tilde{n},\tilde{k},\alpha}| \geq \frac{\binom{\tilde{n}}{\tilde{k}}}{\tilde{n}^2}$; the

existence of these values is guaranteed by Claim 3.5. Let $\Pi_\emptyset = \emptyset$ and consider the Π-matroid instance $I = (\tilde{n}, \tilde{k}, \alpha, M_{\tilde{n},\tilde{k},\alpha}(\Pi_\emptyset))$. Given some string of bits $b \in \{0,1\}^t$, let $Q(b) \subseteq 2^{[\tilde{n}]}$ be the set of containing all $S \subseteq [\tilde{n}]$ queried by $\mathcal{A}$ on the instance I and string b (on the membership oracle of $M_{\tilde{n},\tilde{k},\alpha}(\Pi_\emptyset)$). For clarity, we use $\bar{b}$ for a random string and b for a realization of $\bar{b}$ to a specific string. Note that $Q(b)$ is a set since the algorithm is deterministic for every $b \in \{0,1\}^t$; conversely, $Q(\bar{b})$ is a random set for a random string $\bar{b} \in \{0,1\}^t$. Since the running time of $\mathcal{A}$ on a Π-matroid instance with size $\tilde{n}$ and parameter $\tilde{k}$ is bounded by t, it holds that $|Q(b)| \leq t$ for every $b \in \{0,1\}^t$. Let

$$R(I) = \left\{ S \in \mathcal{F}_{\tilde{n},\tilde{k},\alpha} \;\middle|\; \Pr\left(S \in Q(\bar{b})\right) \geq \frac{1}{2} \right\} \tag{2}$$

be all sets in $\mathcal{F}_{\tilde{n},\tilde{k},\alpha}$ that are queried by $\mathcal{A}$ w.p. at least $\frac{1}{2}$. We now upper bound $|R(I)|$:

$$\begin{aligned}
|R(I)| &= \left| \left\{ S \in \mathcal{F}_{\tilde{n},\tilde{k},\alpha} \;\middle|\; \Pr\left(S \in Q(\bar{b})\right) \geq \frac{1}{2} \right\} \right| \\
&\leq 2 \cdot \sum_{S \in \mathcal{F}_{\tilde{n},\tilde{k},\alpha}} \Pr\left(S \in Q(\bar{b})\right) \\
&= 2 \cdot \sum_{S \in \mathcal{F}_{\tilde{n},\tilde{k},\alpha}} \sum_{b \in \{0,1\}^t} \Pr\left(S \in Q(b)\right) \cdot \Pr(\bar{b} = b)
\end{aligned}$$

Thus, by changing the order of summation

$$\begin{aligned}
|R(I)| &\leq 2 \cdot \sum_{b \in \{0,1\}^t} \Pr(\bar{b} = b) \cdot \sum_{S \in \mathcal{F}_{\tilde{n},\tilde{k},\alpha} \cap Q(b)} \Pr\left(S \in Q(b)\right) \\
&\leq 2 \cdot \sum_{b \in \{0,1\}^t} \Pr(\bar{b} = b) \cdot |Q(b)| \\
&\leq 2t \cdot \sum_{b \in \{0,1\}^t} \Pr(\bar{b} = b) \\
&= 2 \cdot t \\
&< 2c \cdot \frac{\binom{\tilde{n}}{\tilde{k}}}{\tilde{n}^3} \\
&< \frac{\binom{\tilde{n}}{\tilde{k}}}{\tilde{n}^2} \\
&\leq |\mathcal{F}_{\tilde{n},\tilde{k},\alpha}|.
\end{aligned} \tag{3}$$

By (3) there is $S \in F_{\tilde{n},\tilde{k},\alpha} \setminus R(I)$. Consider the Π-matroid instance $I_S = (\tilde{n}, \tilde{k}, \alpha, M_{\tilde{n},\tilde{k},\alpha}(\Pi_S))$ such that $\Pi_S = \{S\}$ and let $B = \{b \in \{0,1\}^t \mid S \notin Q(b)\}$ be all strings for which S is not queried by $\mathcal{A}$. Observe that for all $b \in B$ it holds that the answers to all queries for $T \in Q(b)$ are the same for both oracles (of $M_{\tilde{n},\tilde{k},\alpha}(\Pi_\emptyset)$ and $M_{\tilde{n},\tilde{k},\alpha}(\Pi_S)$); moreover, the decision which set to query next depends only on $b, \tilde{n}, \tilde{k}, \alpha$, and the answers to previous queries. Hence, for all $b \in B$ we have that $\mathcal{A}$ does not query the membership oracle of $M_{\tilde{n},\tilde{k},\alpha}(\Pi_S)$ with S; therefore, $\mathcal{A}$ does not decide I_S correctly for all $b \in B$. We give an illustration in Figure 2. Since $S \notin R(I)$ it holds that $\Pr\left(\bar{b} \in B\right) = \Pr\left(S \notin Q(\bar{b})\right) > \frac{1}{2}$; thus, w.p. at least $\frac{1}{2}$, $\mathcal{A}$ does not decide correctly the instance I_S. This is a contradiction that $\mathcal{A}$ is a randomized algorithm that decides every Π-matroid instance with size $\tilde{n}$ and parameter $\tilde{k}$ in time bounded by t. ■

The following corollary follows directly from Lemma 3.4.

Corollary 3.6. *There is no randomized or deterministic algorithm that decides the* Π*-matroid problem in time* $n^{O(1)}$*, where* n *is the size of the instance.*

4 Tight Lower Bounds using Π-Matroids

In this section we use the hardness of deciding the Π-matroid problem to give tight lower bound on the running time of randomized algorithms for several well studied weighted matroid problems. This includes the budgeted matroid independent set (intersection), exact weight matroid independent set, constrained minimum weight basis of a matroid, and knapsack cover with a a matroid (intersection). We aim to present the results in a compact manner; for this, we rely on the following generic results.

The lower bounds we present in the following are based on the Π-matroid problem and are applied to matroid problems involving weights (e.g., profit, cost, etc.) and a budget (or, a cover demand, a target value). We are able to use in all of the constructions the same family of *standard weight* functions and *standard budgets* in slightly different ways. Specifically, for every $n \in \mathbb{N}_{>0}$ define the standard weight function of n as $w_n : [n] \to \mathbb{N}$ such that for all $i \in [n]$ it holds that

$$w_n(i) = i + 2 \cdot \mathsf{sum}([n]). \tag{4}$$

Moreover, define the standard budget of $n, k, \alpha \in \mathbb{N}_{>0}$ as

$$B_{n,k,\alpha} = \alpha + 2k \cdot \mathsf{sum}([n]). \tag{5}$$

We have the following general results that will be used in the next hardness results. In the next lemma, we give a polynomial bound on the total (standard) weight of any subset of $[n]$.

Lemma 4.1. *For every* $n \in \mathbb{N}_{>0}$ *and* $S \subseteq [n]$*, it holds that* $w_n(S) \in \{0, 1, \ldots, 3 \cdot n^3\}$.

Proof. By (4) we have $w_n(S) \in \mathbb{N}$. In addition, $0 = |\emptyset| \leq |S| \leq \mathsf{sum}(S) \leq w_n(S)$ and

$$w_n(S) \leq n \cdot \max_{i \in [n]} w_n(i) = n \cdot \max_{i \in [n]} i + n \cdot (2 \cdot \mathsf{sum}([n])) \leq n^2 + 2 \cdot n^3 \leq 3 \cdot n^3.$$

■

In the next two results, we show that finding $S \in \mathcal{L}_{n,k,\alpha}(\Pi)$ (for the appropriate parameters) is equivalent to finding $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$ such that $w_n(S) = B_{n,k,\alpha}$. For the first direction,

Lemma 4.2. *For every* $n, k, \alpha \in \mathbb{N}_{>0}, \Pi \subseteq 2^{[n]}$*, and* $S \in \mathcal{L}_{n,k,\alpha}(\Pi)$*, it holds that* $w_n(S) = B_{n,k,\alpha}$.

Proof. Let $S \in \mathcal{L}(\Pi)$. That is, by (1) it holds that $|S| = k$ and $\mathsf{sum}(S) = \alpha$. Then,

$$w_n(S) = \sum_{i \in S} (i + 2 \cdot \mathsf{sum}([n])) = \mathsf{sum}(S) + |S| \cdot 2 \cdot \mathsf{sum}([n]) = \alpha + 2k \cdot \mathsf{sum}([n]) = B_{n,k,\alpha}.$$

■

Now, for the second direction, given $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$ such that $w_n(S) = B_{n,k,\alpha}$ we show that $S \in \mathcal{L}_{n,k,\alpha}(\Pi)$.

Lemma 4.3. *Let* $n, k, \in \mathbb{N}_{>0}, \alpha \in \left[n^2\right], \Pi \subseteq 2^{[n]}$*, and* $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$ *such that* $w_n(S) = B_{n,k,\alpha}$*. Then,* $S \in \mathcal{L}_{n,k,\alpha}(\Pi)$.

Proof. Since $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$, by (1), it is sufficient to show that $|S| = k$ and $\mathsf{sum}(S) = \alpha$. First, assume towards a contradiction that $|S| \neq k$. We consider two cases.

1. $|S| < k$. Then,

$$\begin{aligned} w_n(S) &= \sum_{i \in S} i + |S| \cdot 2 \cdot \mathsf{sum}([n]) \\ &\leq \sum_{i \in S} i + (k-1) \cdot 2 \cdot \mathsf{sum}([n]) \\ &\leq \mathsf{sum}([n]) + (k-1) \cdot 2 \cdot \mathsf{sum}([n]) \\ &< 2k \cdot \mathsf{sum}([n]) \\ &< B_{n,k,\alpha}. \end{aligned} \tag{6}$$

The last inequality holds since $\alpha \geq 1$. By (6) we reach a contradiction that $w_n(S) = B_{n,k,\alpha}$.

2. $|S| > k$. Since $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$, by (1) it holds that $|S| \leq k$ and we reach a contradiction.

By the above we conclude that $|S| = k$. Therefore,

$$\mathsf{sum}(S) = \sum_{i \in S} i = w_n(S) - |S| \cdot 2 \cdot \mathsf{sum}([n]) = w_n(S) - 2k \cdot \mathsf{sum}([n]) = B_{n,k,\alpha} - 2k \cdot \mathsf{sum}([n]) = \alpha.$$

In conclusion, we have $S \in \mathcal{I}(\Pi)$, $|S| = k$ and $\mathsf{sum}(S) = \alpha$. Thus, by (1) it holds that $S \in \mathcal{L}(\Pi)$. ■

The following corollary follows directly from Lemma 4.2 and Lemma 4.3.

Corollary 4.4. *For every $n, k \in \mathbb{N}_{>0}$, $\alpha \in \left[n^2\right]$, $\Pi \subseteq 2^{[n]}$, and $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$ it holds that $S \in \mathcal{L}_{n,k,\alpha}(\Pi)$ if and only if $w_n(S) = B_{n,k,\alpha}$.*

Using Corollary 4.4, we can reduce several weighted matroid problems to the Π-matroid problem. The general idea is, given a Π-matroid instance I, to reduce I to the considered problem using the standard weight and standard budget. The optimal solution for the considered problem is *exactly* $B_{n,k,\alpha}$ (for appropriate parameters) if and only if I has a solution.

4.1 Budgeted Matroid Independent Set

Lemma 4.5. *There is no randomized* FPTAS *for Budgeted Matroid Independent Set.*

Proof. Assume towards a contradiction that there is a randomized FPTAS $\mathcal{A}$ for BM. We use $\mathcal{A}$ to decide the Π-matroid problem. Let $I = (n, k, \alpha, M_{n,k,\alpha}(\Pi))$ be a Π-matroid instance. Consider the following algorithm $\mathcal{B}$ that decides I.

1. Construct the BM instance $U = ([n], \mathcal{I}_{n,k,\alpha}(\Pi), w_n, w_n, B_{n,k,\alpha})$ along with the membership oracle of $M_{n,k,\alpha}(\Pi)$; note that the elements of U are $[n] = \{1, \ldots, n\}$, the independent sets are the independent sets of the Π-matroid $M_{n,k,\alpha}(\Pi)$ (i.e., $\mathcal{I}_{n,k,\alpha}(\Pi)$), the standard weight function w_n is both the profit and the cost function, and the budget is simply the standard budget with parameters n, k, α.
2. Define $\varepsilon = \frac{1}{6 \cdot n^3}$ and compute $\mathcal{A}$ on U and ε; let S be the returned solution.
3. Return that I has a solution if and only if $w_n(S) = B_{n,k,\alpha}$.

Note that U can be constructed from I in time $|I|^{O(1)}$ and that $\frac{1}{\varepsilon} = |I|^{O(1)}$; thus, the running time of $\mathcal{B}$ on I is $|I|^{O(1)}$, since $\mathcal{A}$ is a (randomized) FPTAS. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$; since S is a solution for U it also holds that $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$. Thus, by Corollary 4.4 it holds that I has a solution.

- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\Pi)$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$ implying that T is a solution for U. Therefore, w.p. at least $\frac{1}{2}$:

$$B_{n,k,\alpha} \geq \mathrm{OPT}(U) \geq w_n(S) \geq (1-\varepsilon) \cdot w_n(T) \geq w_n(T) - \frac{1}{6 \cdot n^3} \cdot 3 \cdot n^3 = B_{n,k,\alpha} - \frac{1}{2}. \quad (7)$$

  The first inequality holds since the profit and the cost of each element are the same, and the cost of a solution cannot exceed the budget. The third inequality follows (w.p. at least $\frac{1}{2}$) because S, T are solutions for U, and S is computed by a randomized FPTAS with the parameter ε. The fourth inequality follows from Lemma 4.1 and by the selection of ε. Observe that the value of S is an integer by (4); thus, by (16), it holds that $w_n(S) = B_{n,k,\alpha}$ w.p. at least $\frac{1}{2}$. Then, $\mathcal{B}$ returns that I has a solution w.p. at least $\frac{1}{2}$.

Hence, we decide correctly I w.p. at least $\frac{1}{2}$. It follows that $\mathcal{B}$ is a randomized algorithm which decides the Π-matroid problem, contradicting Lemma 3.4. ■

Corollary 4.6. *There is no randomized* FPTAS *for Budgeted Matroid Intersection.*

4.2 Exact Weight Matroid Independent Set

Lemma 4.7. *For any function $f : \mathbb{N} \to \mathbb{N}$ there is no algorithm that decides* EMI *in time $f(k) \cdot n^{O(1)}$, where n is the number of elements in the given matroid and k is the rank of the matroid.*

Proof. Assume towards a contradiction that there is a function $f : \mathbb{N} \to \mathbb{N}$, a constant q, and an algorithm $\mathcal{A}$ that decides EMI in time $f(k') \cdot n'^q$, where n' is the number of elements in the given matroid and k' is the rank of the matroid. We use $\mathcal{A}$ to decide the Π-matroid problem. By Lemma 3.4, there is a constant c such that there is no randomized algorithm that decides the Π-matroid problem in time $c \cdot \frac{\binom{n'}{k}}{n'^3}$. Let $I = (n, k, \alpha, M_{n,k,\alpha}(\Pi))$ be a Π-matroid instance such that $6 \cdot n^3 + f(k) \cdot n^q < c \cdot \frac{\binom{n}{k}}{n^3}$. Consider the following algorithm $\mathcal{B}$ that decides I.

1. Construct the EMI instance $U = ([n], \mathcal{I}_{n,k,\alpha}(\Pi), w_n, k, B_{n,k,\alpha})$ along with the membership oracle of $M_{n,k,\alpha}(\Pi)$; note that the elements of U are $[n] = \{1, \ldots, n\}$, the independent sets are the independent sets of the Π-matroid $M_{n,k,\alpha}(\Pi)$ (i.e., $\mathcal{I}_{n,k,\alpha}(\Pi)$), the standard weight function w_n is the weight function, the target cardinality is k, and the target value is the standard budget with parameters n, k, α. We say that $T \subseteq [n]$ is a solution for U if $T \in \mathcal{I}_{n,k,\alpha}$ and $w_n(T) = B_{n,k,\alpha}$.
2. Compute $\mathcal{A}$ on U.
3. Return that I has a solution if and only if $\mathcal{A}$ returns that U has a solution.

Note that U can be constructed naively from I in time $n + n \cdot (2n^2 + n) + 2n \cdot k + n^2 \leq 6 \cdot n^3$ (recall that the weights are encoded in unary, $k \leq n$, and that in the oracle model we do not explicitly encode the independent sets of the matroid); thus, the running time of $\mathcal{B}$ on I is bounded by $6 \cdot n^3 + f(k) \cdot n^q$. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$ and $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$. Thus, by Corollary 4.4 it holds that I has a solution.
- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\Pi)$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$ implying that T is a solution for U. Therefore, w.p. at least $\frac{1}{2}$ it holds that $\mathcal{A}$ returns that U has a solution; thus, w.p. at least $\frac{1}{2}$ it holds that $\mathcal{B}$ returns that I has a solution.

Hence, we decide correctly I w.p. at least $\frac{1}{2}$. It follows that $\mathcal{B}$ is a randomized algorithm which decides the Π-matroid problem in time $6 \cdot n^3 + f(k) \cdot n^q < c \cdot \frac{\binom{n}{k}}{n^3}$, contradicting Lemma 3.4. ■

4.3 Constrained Minimum Basis of a Matroid

In contrary to BM, EMI, and KCM, the reduction from Π-matroid to CMB requires different *cost* function. Specifically, for every $n \in \mathbb{N}_{>0}$ define the standard cost function of n as $c_n : [n] \to \mathbb{N}$ such that for all $i \in [n]$ it holds that

$$c_n(i) = w_n(i) - 2 \cdot i = 2 \cdot \mathsf{sum}([n]) - i. \tag{8}$$

Moreover, define the standard bound of $n, k, \alpha \in \mathbb{N}_{>0}$ as

$$L_{n,k,\alpha} = B_{n,k,\alpha} - 2 \cdot \alpha = 2k \cdot \mathsf{sum}([n]) - \alpha. \tag{9}$$

The above defined cost function is useful for the proof of hardness of CMB. We first use the following auxiliary results.

Lemma 4.8. *For every $n, k, \alpha \in \mathbb{N}_{>0}$, and $S \subseteq [n], |S| = k$ it holds that $w_n(S) \geq B_{n,k,\alpha}$ if and only if $c_n(S) \leq L_{n,k,\alpha}$.*

Proof. If $w_n(S) \geq B_{n,k,\alpha}$, then by (4), (5), and since $|S| = k$ we have:

$$2 \cdot \mathsf{sum}([n]) \cdot k + \alpha = B_{n,k,\alpha} \leq w_n(S) = 2 \cdot \mathsf{sum}([n]) \cdot k + \mathsf{sum}(S). \tag{10}$$

By (10) we conclude that $\mathsf{sum}(S) \geq \alpha$. Therefore,

$$c_n(S) = 2 \cdot \mathsf{sum}([n]) \cdot k - \mathsf{sum}(S) \leq 2 \cdot \mathsf{sum}([n]) \cdot k - \alpha = L_{n,k,\alpha}.$$

For the second direction, if $c_n(S) \leq L_{n,k,\alpha}$, then by (8), (9), and since $|S| = k$ it holds that:

$$2 \cdot \mathsf{sum}([n]) \cdot k - \alpha = L_{n,k,\alpha} \geq c_n(S) = 2 \cdot \mathsf{sum}([n]) \cdot k - \mathsf{sum}(S). \tag{11}$$

By (11) we conclude that $\mathsf{sum}(S) \geq \alpha$. Then,

$$w_n(S) = 2 \cdot \mathsf{sum}([n]) \cdot k + \mathsf{sum}(S) \geq 2 \cdot \mathsf{sum}([n]) \cdot k + \alpha = B_{n,k,\alpha}.$$

■

Lemma 4.9. *There is no randomized* FPTAS *for Constrained Minimum Basis of a Matroid.*

Proof. Assume towards a contradiction that there is a randomized FPTAS $\mathcal{A}$ for CMB. We use $\mathcal{A}$ to decide the Π-matroid problem. Let $I = (n, k, \alpha, M_{n,k,\alpha}(\Pi))$ be a Π-matroid instance. Consider the following algorithm $\mathcal{B}$ that decides I.

1. Construct the CMB instance $U = ([n], \mathcal{I}_{n,k,\alpha}(\Pi), c_n, w_n, B_{n,k,\alpha})$ along with the membership oracle of $M_{n,k,\alpha}(\Pi)$; note that the elements of U are $[n] = \{1, \ldots, n\}$, the independent sets are the independent sets of the Π-matroid $M_{n,k,\alpha}(\Pi)$ (i.e., $\mathcal{I}_{n,k,\alpha}(\Pi)$), the cost is the standard cost function, the weight is the standard weight function w_n, and the bound is the standard bound with parameters n, k, α.
2. Define $\varepsilon = \frac{1}{6 \cdot n^3}$ and compute $\mathcal{A}$ on U and ε; let S be the returned solution.
3. Return that I has a solution if and only if $w_n(S) = B_{n,k,\alpha}$.

Note that U can be constructed from I in time $|I|^{O(1)}$ and that $\frac{1}{\varepsilon} = |I|^{O(1)}$; thus, the running time of $\mathcal{B}$ on I is $|I|^{O(1)}$, since $\mathcal{A}$ is a (randomized) FPTAS. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$; since S is a solution for U it also holds that $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$. Thus, by Corollary 4.4 it holds that I has a solution.

- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\Pi)$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$, which implies by Lemma 4.8 that $c_n(T) \leq L_{n,k,\alpha}$. Then, since T is a basis of $M_{n,k,\alpha}(\Pi)$ ($T \in \mathcal{L}_{n,k,\alpha}(\Pi)$), it follows that T is a solution for U. Hence, w.p. at least $\frac{1}{2}$:

$$B_{n,k,\alpha} \leq \mathrm{OPT}(U) \leq w_n(S) \leq (1+\varepsilon) \cdot w_n(T) \leq w_n(T) + \frac{1}{6 \cdot n^3} \cdot 3 \cdot n^3 = B_{n,k,\alpha} + \frac{1}{2}. \quad (12)$$

  The first inequality follows from Lemma 4.8 (note that U has a solution since T is a solution). The third inequality follows (w.p. at least $\frac{1}{2}$) because S, T are solutions for U, and S is computed by a randomized FPTAS with the parameter ε. The fourth inequality follows from Lemma 4.1 and by the selection of ε. Observe that the value of S is an integer by (4); thus, by (15), it holds that $w_n(S) = B_{n,k,\alpha}$ w.p. at least $\frac{1}{2}$. Then, $\mathcal{B}$ returns that I has a solution w.p. at least $\frac{1}{2}$.

Hence, we decide correctly I w.p. at least $\frac{1}{2}$. It follows that $\mathcal{B}$ is a randomized algorithm which decides the Π-matroid problem, contradicting Lemma 3.4.

■

4.4 Knapsack Cover with a Matroid

Proof. Assume towards a contradiction that there is a randomized FPTAS $\mathcal{A}$ for KCM. We use $\mathcal{A}$ to decide the Π-matroid problem. Let $I = (n, k, \alpha, M_{n,k,\alpha}(\Pi))$ be a Π-matroid instance. Consider the following algorithm $\mathcal{B}$ that decides I.

1. Construct the KCM instance $U = ([n], \mathcal{I}_{n,k,\alpha}(\Pi), w_n, w_n, B_{n,k,\alpha})$ along with the membership oracle of $M_{n,k,\alpha}(\Pi)$; note that the elements of U are $[n] = \{1, \ldots, n\}$, the independent sets are the independent sets of the Π-matroid $M_{n,k,\alpha}(\Pi)$ (i.e., $\mathcal{I}_{n,k,\alpha}(\Pi)$), the standard weight function w_n is both the cost and the size function, and the demand is simply the standard budget with parameters n, k, α.

2. Define $\varepsilon = \frac{1}{6 \cdot n^3}$ and compute $\mathcal{A}$ on U and ε; let S be the returned solution.

3. Return that I has a solution if and only if $w_n(S) = B_{n,k,\alpha}$.

Note that U can be constructed from I in time $|I|^{O(1)}$ and that $\frac{1}{\varepsilon} = |I|^{O(1)}$; thus, the running time of $\mathcal{B}$ on I is $|I|^{O(1)}$, since $\mathcal{A}$ is a (randomized) FPTAS. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$; since S is a solution for U it also holds that $S \in \mathcal{I}_{n,k,\alpha}(\Pi)$. Thus, by Corollary 4.4 it holds that I has a solution.

- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\Pi)$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$ implying that T is a solution for U. Therefore, w.p. at least $\frac{1}{2}$:

$$B_{n,k,\alpha} \leq \mathrm{OPT}(U) \leq w_n(S) \leq (1+\varepsilon) \cdot w_n(T) \leq w_n(T) + \frac{1}{6 \cdot n^3} \cdot 3 \cdot n^3 = B_{n,k,\alpha} + \frac{1}{2}. \quad (13)$$

  The first inequality holds since the size and the cost of each element are the same, and the size of a solution must be higher or equal from the demand. The third inequality follows (w.p.

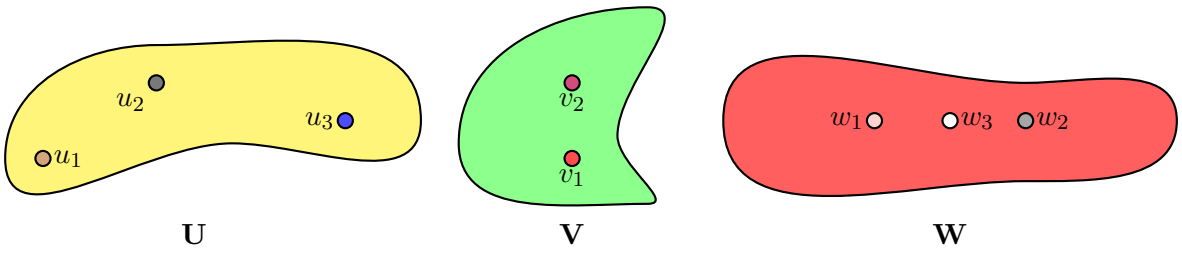

Figure 3: An example of a *partition matroid* $(E, \mathcal{I})$ - a matroid that can be efficiently encoded. The ground set is $E = \{u_1, u_2, u_3, v_1, v_2, w_1, w_2, w_3\}$, partitioned into three sets U, V, W. The independent sets are all subsets of E containing at most one element from U, V, and W; that is, $\mathcal{I} = \{S \subseteq E \mid \forall X \in \{U, V, W\} : |S \cap X| \leq 1\}$. A simple efficient encoding of $(E, \mathcal{I})$ is $I = (E, U, V, W)$. Membership can be decided efficiently given I by checking the feasibility of each set U, V, W.

at least $\frac{1}{2}$) because S, T are solutions for U, and S is computed by a randomized FPTAS with the parameter ε. The fourth inequality follows from Lemma 4.1 and by the selection of ε. Observe that the value of S is an integer by (4); thus, by (16), it holds that $w_n(S) = B_{n,k,\alpha}$ w.p. at least $\frac{1}{2}$. Then, $\mathcal{B}$ returns that I has a solution w.p. at least $\frac{1}{2}$.

Hence, we decide correctly I w.p. at least $\frac{1}{2}$. It follows that $\mathcal{B}$ is a randomized algorithm which decides the Π-matroid problem, contradicting Lemma 3.4.

■

5 The SAT-Matroid Family

The hardness result described in Section 3 shows unconditionally that the Π-matroid problem is very hard to solve. Nonetheless, this hardness result considers matroids with general membership oracles and does not give a lower bound for matroids that can be explicitly encoded. This is particularly important as matroids are often used to describe problems with an explicit encoding. This include partition matroids, graphic matroids, linear matroids, etc.[2] We formally define what we mean by explicitly encoded and define an encoding of matroids that do not rely on a membership oracle.

Definition 5.1. *A function* $f : \{0,1\}^* \to \mathbb{N} \times 2^{\mathbb{N}}$ *is called* matroid encoder *if for every* $I \in \{0,1\}^*$ *it holds that* $f(I) = \left(E_{f(I)}, \mathcal{I}_{f(I)}\right)$ *is a matroid such that the following holds.*

1. *There is an algorithm that given* $I \in \{0,1\}^*$ *returns* $E_{f(I)}$ *in time* $|I|^{O(1)}$.
2. *There is an algorithm that given* $I \in \{0,1\}^*$ *and* $S \subseteq E_{f(I)}$ *decides if* $S \in \mathcal{I}_{f(I)}$ *in time* $|I|^{O(1)}$.

Note that every matroid has a matroid encoder; however, the encoding may be inefficient: a matroid $(E, \mathcal{I})$ encoded by $I \in \{0,1\}^*$ using a matroid encoder $f : \{0,1\}^* \to \mathbb{N} \times 2^{\mathbb{N}}$ (i.e., $f(I) = (E, \mathcal{I})$) may satisfy that $|I| = \Omega\left(2^{|E|}\right)$. This difficulty in general matroid representation has lead to the common use of membership oracles, which allow to analyze the complexity of algorithms for matroid problems considering only $|E|$ (and not the encoding of $\mathcal{I}$). This could

[2] See, e.g., [?] for a survey on various families of matroids.

have suggested that the hardness of the Π-matroid problem is rooted in the hardness that the membership oracle model manifest. Yet, there are matroids with a very efficient encoding. For examples, a *uniform matroid* $(E, \mathcal{I})$ such that $\mathcal{I} = \{S \subseteq E \mid |S| \leq k\}$ can be efficiently encoded using $I = (E, k) \in \{0,1\}^*$; clearly, deciding membership of some $S \subseteq E$ only depends on $|S|$ and k. Another example is given in Figure 3. We show that there are matroids that can be efficiently encoded (i.e., $|I| = |E|^{O(1)}$) but their corresponding Π-matroid problem is NP-Hard.

We use a reduction from the known SAT problem to a variation of the Π-matroid problem where the matroid is explicitly given by a matroid encoder. The main idea is to use an efficiently encodable member of the Π-matroid family, in which the secret set Π contains all feasible solutions of a given SAT instance. In a SAT instance $A = (V, \bar{V}, \mathcal{C})$ with $n \in \mathbb{N}$ variables (in a slightly simplified notation), we are given a set $V = \{v_1, \ldots, v_n\}$ of variables, their negations $\bar{V} = \{\bar{v}_1, \ldots, \bar{v}_n\}$, and a set $\mathcal{C} \subseteq 2^{V \cup \bar{V}}$ of clauses. The goal is to decide if there is a set $S \subseteq [n]$ satisfying that for all $C \in \mathcal{C}$ there is $i \in [n]$ such that one of the following holds.

- $v_i \in C$ and $i \in S$.
- $\bar{v}_i \in C$ and $i \notin S$.

Such a set S is called a *solution* of A and let $\mathcal{S}(A)$ be the set of solutions of a SAT instance A. In addition, let $n(A) = n$ be the number of variables in the instance A. In the *parametrized* variant of SAT, in addition to a SAT instance A we are given a parameter k; the goal is to decide if there is a solution S of A such that $|S| = k$. We use (A, k) to denote parametrized SAT instances, where A is a SAT instance and k is the parameter; also, let $\mathcal{S}_k(A)$ be all solutions for A of cardinality k. The family of SAT-matroids is the subfamily of Π-matroids where $\Pi = \mathcal{S}(A)$ for some SAT instance A.[3] Formally,

Definition 5.2. *Let A be a* SAT *instance, $k \in [n(A)]$, and $\alpha \in \left[n(A)^2\right]$. Define the* **SAT-matroid** *on A, k, α as $M_{n(A),k,\alpha}(\mathcal{S}(A)) = \left([n], \mathcal{I}_{n(A),k,\alpha}(\mathcal{S}(A))\right)$.*

We show below that the SAT-matroids can be efficiently encoded without the use of membership oracles. For every $I \in \{0,1\}^*$, we interpret I as $I = (A, k, \alpha)$, where A is a SAT instance, $k \in [n(A)]$, and $\alpha \in \left[n(A)^2\right]$; w.l.o.g., we may assume that every $I \in \{0,1\}^*$ can be interpreted in that manner and that $n(A) \leq |A|$. Additionally, for simplicity we use $n_I = n(A), k_I = k, \alpha_I = \alpha$.

Definition 5.3. *Define the* **SAT-encoder** *function $f_{\text{SAT}} : \{0,1\}^* \to \mathbb{N} \times 2^{\mathbb{N}}$ such that for all $I \in \{0,1\}^*$ it holds that $f_{\text{SAT}}(I) = M_{n_I,k_I,\alpha_I}(\mathcal{S}(A_I))$.*

In the next result, we show that the SAT-encoder f_{SAT} is not only a matroid encoder, but it does so with an efficient encoding. This is crucial of reproducing a similar variant of the Π-matroid problem without a membership oracle involved.

Lemma 5.4. *The* SAT-encoder *f_{SAT} is a matroid encoder such that for every $I \in \{0,1\}^*$ it holds that $|I| = |A_I|^{O(1)}$.*

Proof. We first show that f_{SAT} is a matroid encoder. For all $I \in \{0,1\}^*$ it holds that $\mathcal{S}(A_I) \subseteq 2^{[n_I]}$; thus, by Lemma 3.2 it holds that f_{SAT} returns a matroid. We now prove that f_{SAT} satisfies the two conditions of Definition 5.1.

1. For the first condition of Definition 5.1, given $I \in \{0,1\}^*$ we can compute $E_{f_{\text{SAT}}(I)} = [n_I]$ in time $|I|^{O(1)}$ by identifying the number of variables (i.e., n_I) in A.

[3] Recall that Π-matroids are defined in Section 3.

2. For the second condition of Definition 5.1, let $I \in \{0,1\}^*$; we show that we can decide membership in the matroid in efficient time. For every $S \subseteq [n_I]$ we can compute $|S|$ and $\mathsf{sum}(S)$ in time $|I|^{O(1)}$ as I is encoded in unary. Thus, validating if $S \in \mathcal{J}_{n_I,k_I}$ or $S \in \mathcal{K}_{n_I,k_I,\alpha_I}$ can be computed in time $|I|^{O(1)}$ regardless of $\mathcal{S}(A_I)$. In addition, deciding if $S \in \mathcal{S}(A_I)$ can be computed in time $|A_I|^{O(1)} = |I|^{O(1)}$ by validating that all clauses are satisfied. Thus, by (1) we can decide if $S \in \mathcal{L}_{n_I,k_I,\alpha_I}(\mathcal{S}(A_I))$. Overall, deciding whether $S \in \mathcal{J}_{n_I,k_I}$, $S \in \mathcal{K}_{n_I,k_I,\alpha_I}$, or $S \in \mathcal{L}_{n_I,k_I,\alpha_I}(\mathcal{S}(A_I))$ can be computed in time $|I|^{O(1)}$. By the above and (1), we can compute in time $|I|^{O(1)}$ correctly if a set S belongs to $\mathcal{I}_{f_{\mathrm{SAT}}(I)} = \mathcal{I}_{n_I,k_I,\alpha_I}(\mathcal{S}(A_I))$.

Thus, f_{SAT} is a matroid encoder by Definition 5.1. To conclude, observe that for all $I \in \{0,1\}^*$ it holds that $k_I, \alpha_I = {n_I}^{O(1)}$ and $n_I \leq |A|$ by the definition of our encoding; hence, $|I| = |A|^{O(1)}$. ■

The above lemma guarantees the efficient representation of SAT-matroids. This makes the following decision problem non-trivial to solve in polynomial time (in the encoding size of the instance). In fact, subsequently we show it is strongly NP-hard. For the following, recall that $\mathcal{L}_{n,k,\alpha}(\Pi)$ is defined in (1).

Definition 5.5. *The* **SAT-matroid problem:** *Given $I \in \{0,1\}^*$ decide if $\mathcal{L}_{n_I,k_I,\alpha_I}(\mathcal{S}(A_I)) \neq \emptyset$, where I is interpreted as A_I, k_I, and α_I.*

For a SAT-matroid instance $I \in \{0,1\}^*$, A *solution* of I is a set $S \in \mathcal{L}_{n_I,k_I,\alpha_I}(\mathcal{S}(A_I))$.

Lemma 5.6. *For every computable function $g, f : \mathbb{N} \to \mathbb{N}$, if there is an algorithm that decides every* SAT*-matroid instance $I \in \{0,1\}^*$ in time $O(f(k_I) \cdot g(|I|))$, then there is an algorithm that decides every parametrized* SAT *instance (A,k) in time $O\left(f(k) \cdot |A|^2 \cdot g\left(O\left(|A|^2\right)\right)\right)$.*

Proof. Let $\mathcal{A}$ be an algorithm for SAT-matroid that decides every SAT-matroid instance $I \in \{0,1\}^*$ in time $O(f(k_I) \cdot g(|I|))$, where $f, g : \mathbb{N} \to \mathbb{N}$ are computable functions. Let (A,k) be a parametrized SAT instance and let $n = n(A)$. We describe the following algorithm which decides (A,k).

1. For every $\alpha \in \left[n^2\right]$, compute algorithm $\mathcal{A}$ on the SAT-matroid $I_{A,k,\alpha} = (A,k,\alpha)$.
2. If there is $\alpha \in \left[n^2\right]$ such that algorithm $\mathcal{A}$ on instance $I_{A,k,\alpha}$ decides that $\mathcal{L}_{n,k,\alpha}(\mathcal{S}(A)) \neq \emptyset$, we return that (A,k) has a solution; otherwise, we return that (A,k) does not have a solution.

For every $\alpha \in \left[n^2\right]$, observe that a solution $S \in \mathcal{L}_{n,k,\alpha}(\mathcal{S}(A))$ for the SAT-matroid instance $I_{A,k,\alpha}$ is in particular a solution $S \in \mathcal{S}_k(A)$ for (A,k). Thus, to show that the above algorithm correctly decides A (and the parametrized SAT problem), consider the following cases.

- The instance (A,k) has a solution $S \in \mathcal{S}_k(A)$. Let $\alpha = \mathsf{sum}(S)$. By (1) it follows that the SAT-matroid instance $I_{A,k,\alpha}$ has a solution, which is S. Therefore, when algorithm $\mathcal{A}$ decides instance $I_{A,k,\alpha}$, it must decide that $\mathcal{L}_{n,k,\alpha}(\mathcal{S}(A)) \neq \emptyset$
- The instance (A,k) does not have a solution, i.e., $\mathcal{S}_k(A) = \emptyset$. Assume towards a contradiction that there is $\alpha \in \left[n^2\right]$ such that there is a solution S for $I_{A,k,\alpha}$. Thus, $S \in \mathcal{L}_{n,k,\alpha}(\mathcal{S}(A))$ by Definition 5.5; therefore, since $\mathcal{L}_{n,k,\alpha}(\mathcal{S}(A)) \subseteq \mathcal{S}_k(A)$ it follows that $S \in \mathcal{S}_k(A)$ in contradiction that $\mathcal{S}_k(A) = \emptyset$.

Assume that $n \leq |A|$ as A contains n distinct variables; then, observe that $|I_{A,k,\alpha}| = |A| + k + \alpha \leq |A| + n + n^2 = O(|A|^2)$. Thus, the number of iterations (i.e., distinct options for the parameters k, α) is bounded by $O(|A|^2)$, and the running time of each iteration is bounded by $O(f(k) \cdot g(O(|A|^2)))$. Therefore, our algorithm decides parametrized SAT in time $O(f(k) \cdot |A|^2 \cdot g(O(|A|^2)))$. ■

The following corollary is a consequence of Lemma 5.6.

Lemma 5.7. *Unless* P=NP*, there is no polynomial time algorithm that decides the* SAT*-matroid problem in time* $|I|^{O(1)}$*, where* I *is the given* SAT*-matroid instance.*

Proof. Assume that P $\neq NP$ and assume towards a contradiction that there is an algorithm $\mathcal{A}$ that decides the SAT-matroid problem in time $|I|^{O(1)}$. Therefore, by Lemma 5.6 there is an algorithm $\mathcal{B}$ that decides parametrized SAT in time $|I|^{O(1)}$. Let A be a SAT instance and let $n = n(A)$; consider the following algorithm $\mathcal{C}$ that decides SAT based on $\mathcal{B}$.

1. Check if $\emptyset \in \mathcal{S}(A)$.
2. For every $k \in [n]$ compute $\mathcal{B}$ on (A, k).
3. If there is $k \in [n]$ such that algorithm $\mathcal{B}$ on instance (A, k) decides that (A, k) has a solution, or that $\emptyset \in \mathcal{S}(A)$, we return that A has a solution; otherwise, we return that A does not have a solution.

Since $\mathcal{B}$ runs in time $|I|^{O(1)}$ and $\mathcal{C}$ computes $\mathcal{A}$ at most $n \leq |I|$ times, the running time of $\mathcal{C}$ is also $|I|^{O(1)}$. For the correctness, consider the following cases.

- If A has a solution. Then, there is $S \in \mathcal{S}(A)$; therefore, either $S = \emptyset$, or, for $k = |S|$ it holds that (A, k) has a solution. Thus, $\mathcal{C}$ returns that A has a solution.
- If $\mathcal{C}$ returns that A has a solution; then, either $\emptyset \in \mathcal{S}(A)$, or, by the definition of $\mathcal{B}$, there is $k \in [n]$ such that (A, k) has a solution, which is $S \in \mathcal{S}_k(A)$. Hence, $\mathcal{S}(A) \neq \emptyset$ and A has a solution.

Hence, $\mathcal{C}$ a polynomial time algorithm that decides decides SAT. Since SAT is well-known to be NP-Hard [7], we reach a contradiction and the proof follows. ■

By Lemma 5.7 we also have the following parametrized hardness result for SAT-matroid.

Lemma 5.8. *Unless* W[1]=FPT*, for every computable function* $f : \mathbb{N} \to \mathbb{N}$ *there is no algorithm that decides the* SAT*-matroid problem in time* $f(k_I) \cdot n_I^{O(1)}$*, where* I *is the* SAT*-matroid instance.*

Proof. s ■

6 Tight Lower Bounds using SAT-Matroids

6.1 Budgeted Matroid Independent Set

We define the variation of BM where the given matroid is explicitly encoded (see Definition 5.1).

f-Budgeted Matroid Independent Set (f-BM)

Encoder	$f : \{0,1\}^* \to \mathbb{N} \times 2^{\mathbb{N}}$ is a matroid encoder.
Instance	$U = (I, c, p, B)$, where $I \in \{0,1\}^*$, $c : E_{f(I)} \to \mathbb{N}$ is a cost function, $p : E_{f(I)} \to \mathbb{N}$ is a profit function and $B \in \mathbb{N}$ is a budget.
Solution	$S \in \mathcal{I}_{f(I)}$ such that $c(S) \leq B$.
Objective	Find a solution S of U with maximum profit $p(S)$.

Lemma 6.1. *Unless* P=NP*, there is no* FPTAS *for* f_{SAT}-BM.

Proof. Assume that P≠NP and assume towards a contradiction that there is an FPTAS $\mathcal{A}$ for f_{SAT}-BM. We use $\mathcal{A}$ to decide the Π-matroid problem. Let $I \in \{0,1\}^*$ be a SAT-matroid instance. Consider the following algorithm $\mathcal{B}$ that decides I. For simplicity, let $A = A_I, n = n(A), k = k_I$, and $\alpha = \alpha_I$.

1. Construct the f_{SAT}-BM instance $U = (I, w_n, w_n, B_{n,k,\alpha})$. Note that $f_{\text{SAT}}(I) = M_{n,k,\alpha}(\mathcal{S}(A))$ by Definition 5.3.
2. Define $\varepsilon = \frac{1}{6 \cdot n^3}$ and compute $\mathcal{A}$ on U and ε; let S be the returned solution.
3. Return that I has a solution if and only if $w_n(S) = B_{n,k,\alpha}$.

By Lemma 5.4, (4), and (5) it follows that U can be constructed from I in time $|I|^{O(1)}$. In addition, $\frac{1}{\varepsilon} = n^{O(1)} = |I|^{O(1)}$; thus, the running time of $\mathcal{B}$ on I is $|I|^{O(1)}$, since $\mathcal{A}$ is an FPTAS. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$; since S is a solution for U it also holds that $S \in \mathcal{I}_{f_{\text{SAT}}(I)} = \mathcal{I}_{n,k,\alpha}(\mathcal{S}(A))$. Thus, by Corollary 4.4 it holds that I has a solution.
- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\mathcal{S}(A))$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$ implying that T is a solution for U. Therefore,

$$B_{n,k,\alpha} \geq \text{OPT}(U) \geq w_n(S) \geq (1-\varepsilon) \cdot w_n(T) \geq w_n(T) - \frac{1}{6 \cdot n^3} \cdot 3 \cdot n^3 = B_{n,k,\alpha} - \frac{1}{2}. \quad (14)$$

  The first inequality holds since the profit and the cost of each element are the same, and the cost of a solution cannot exceed the budget. The third inequality follows because S, T are solutions for U, and S is computed by an FPTAS with the parameter ε. The fourth inequality follows from Lemma 4.1 and by the selection of ε. Observe that the value of S is an integer by (4); thus, by (16), it holds that $w_n(S) = B_{n,k,\alpha}$. Then, $\mathcal{B}$ returns that I has a solution.

Hence, we decide correctly I. It follows that $\mathcal{B}$ is a polynomial algorithm that decides the SAT-matroid problem, contradicting Lemma 5.7. ■

6.2 Exact Weight Matroid Independent Set

f-Exact Matroid Independent Set (f-EMI)

Encoder	$f : \{0,1\}^* \to \mathbb{N} \times 2^{\mathbb{N}}$ is a matroid encoder.
Instance	$U = (I, w, k, T)$, where $I \in \{0,1\}^*$, $w : E_{f(I)} \to \mathbb{N}$ gives the elements weights, $k \in \mathbb{N}$ is the target cardinality, and $T \in \mathbb{N}$ is the target weight.
Objective	Decide if there is $S \in \mathcal{I}_{f(I)}$ such that $\lvert S\rvert = k$ and $w(S) = T$.

Lemma 6.2. *Unless* W[1]=FPT*, for any function $f : \mathbb{N} \to \mathbb{N}$ there is no algorithm that decides* EMI *in time $f(k) \cdot n^{O(1)}$, where n is the number of elements in the given matroid and k is the rank of the matroid.*

Proof. Assume that W[1]≠FPT and assume towards a contradiction that there is a function $f : \mathbb{N} \to \mathbb{N}$, a constant q, and an algorithm $\mathcal{A}$ that decides f_{SAT}-EMI in time $f(k') \cdot n'^q$, where n' is the number of elements in the given matroid and k' is the rank of the matroid. We use $\mathcal{A}$ to decide the SAT-matroid problem. Let $I \in \{0,1\}^*$ be a SAT-matroid instance. Consider the following algorithm $\mathcal{B}$ that decides I. For simplicity, let $A = A_I, n = n(A), k = k_I$, and $\alpha = \alpha_I$. Assume that $k \leq n$ or the instance can be trivially decided.

1. Construct the f_{SAT}-EMI instance $U = (I, w_n, k, B_{n,k,\alpha})$. Note that $f_{\text{SAT}}(I) = M_{n,k,\alpha}(\mathcal{S}(A))$ by Definition 5.3.

2. Compute $\mathcal{A}$ on U.

3. Return that I has a solution if and only if $\mathcal{A}$ returns that U has a solution.

Note that U can be constructed naively from I in time $n + n \cdot (2n^2 + n) + 2n \cdot k + n^2 \leq 6 \cdot n^3$ (recall that the weights are encoded in unary and that $k \leq n$); thus, the running time of $\mathcal{B}$ on I is bounded by $6 \cdot n^3 + f(k) \cdot n^q = f(k) \cdot n^{O(1)}$. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$ and $S \in \mathcal{I}_{f_{\text{SAT}}(I)} = \mathcal{I}_{n,k,\alpha}(\mathcal{S}(A))$. Thus, by Corollary 4.4 it holds that I has a solution.

- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\mathcal{S}(A))$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$ implying that T is a solution for U. Therefore, $\mathcal{A}$ returns that U has a solution; thus, $\mathcal{B}$ returns that I has a solution.

Hence, we decide I correctly. It follows that $\mathcal{B}$ is an algorithm which decides the SAT-matroid problem in time $f(k) \cdot n^{O(1)}$, contradicting Lemma 3.4. ■

6.3 Constrained minimum Basis of a Matroid

	f-Constrained minimum Basis of a Matroid (f-CMB)
Encoder	$f : \{0,1\}^* \to \mathbb{N} \times 2^{\mathbb{N}}$ is a matroid encoder.
Instance	$U = (I, c, w, B)$, where $I \in \{0,1\}^*$, $c : E_{f(I)} \to \mathbb{N}$ is a cost function, $w : E_{f(I)} \to \mathbb{N}$ is a weight function and $B \in \mathbb{N}$ is a budget.
Solution	A basis $S \in \mathcal{I}_{f(I)}, \lvert S\rvert = \max_{S' \in \mathcal{I}_{f(I)}} \lvert S'\rvert$ such that $c(S) \leq B$.
Objective	Find a solution S of U with minimum weight $w(S)$.

Lemma 6.3. *Unless* P=NP*, there is no* FPTAS *for* f_{SAT}-CMB.

Proof. Assume that P≠NP and assume towards a contradiction that there is an FPTAS $\mathcal{A}$ for f_{SAT}-CMB. We use $\mathcal{A}$ to decide the SAT-matroid problem. Let $I \in \{0,1\}^*$ be a SAT-matroid instance. Consider the following algorithm $\mathcal{B}$ that decides I. For simplicity, let $A = A_I, n = n(A), k = k_I$, and $\alpha = \alpha_I$.

1. Construct the f_{SAT}-CMB instance $U = (I, c_n, w_n, B_{n,k,\alpha})$. Note that $f_{\text{SAT}}(I) = M_{n,k,\alpha}(\mathcal{S}(A))$ by Definition 5.3.

2. Define $\varepsilon = \frac{1}{6 \cdot n^3}$ and compute $\mathcal{A}$ on U and ε; let S be the returned solution.

3. Return that I has a solution if and only if $w_n(S) = B_{n,k,\alpha}$.

By Lemma 5.4, (4), and (5) it follows that U can be constructed from I in time $\lvert I\rvert^{O(1)}$. In addition, $\frac{1}{\varepsilon} = n^{O(1)} = \lvert I\rvert^{O(1)}$; thus, the running time of $\mathcal{B}$ on I is $\lvert I\rvert^{O(1)}$, since $\mathcal{A}$ is an FPTAS. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$; since S is a solution for U it also holds that $S \in \mathcal{I}_{f_{\text{SAT}}(I)} = \mathcal{I}_{n,k,\alpha}(\mathcal{S}(A))$. Thus, by Corollary 4.4 it holds that I has a solution.

- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\mathcal{S}(A))$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$, which implies by Lemma 4.8 that $c_n(T) \leq L_{n,k,\alpha}$. Then, since T is a basis of $M_{n,k,\alpha}(\Pi)$ (recall that $T \in \mathcal{L}_{n,k,\alpha}(\mathcal{S}(A))$), it follows that T is a solution for U. Hence,

$$B_{n,k,\alpha} \leq \mathrm{OPT}(U) \leq w_n(S) \leq (1+\varepsilon) \cdot w_n(T) \leq w_n(T) + \frac{1}{6 \cdot n^3} \cdot 3 \cdot n^3 = B_{n,k,\alpha} + \frac{1}{2}. \quad (15)$$

  The first inequality follows from Lemma 4.8 (note that U has a solution since T is a solution). The third inequality follows because S, T are solutions for U, and S is computed by an FPTAS with the parameter ε. The fourth inequality follows from Lemma 4.1 and by the selection of ε. Observe that the value of S is an integer by (4); thus, by (15), it holds that $w_n(S) = B_{n,k,\alpha}$. Then, $\mathcal{B}$ returns that I has a solution.

Hence, we decide I correctly. It follows that $\mathcal{B}$ is a polynomial algorithm which decides the SAT-matroid problem, contradicting Lemma 5.7. ■

6.4 Knapsack Cover with a Matroid

f-Knapsack Cover with a Matroid (f-KCM)

Encoder	$f : \{0,1\}^* \to \mathbb{N} \times 2^{\mathbb{N}}$ is a matroid encoder.
Instance	$U = (I, c, d, D)$, where $I \in \{0,1\}^*$, $c : E_{f(I)} \to \mathbb{N}$ is a cost function, $d : E_{f(I)} \to \mathbb{N}$ is a size function, and $B \in \mathbb{N}$ is a demand.
Solution	$S \in \mathcal{I}_{f(I)}$ such that $d(S) \geq D$.
Objective	Find a solution S of U with minimum cost $c(S)$.

Lemma 6.4. *Unless* P=NP, *there is no* FPTAS *for* f_{SAT}-KCM.

Proof. Assume that P≠NP and assume towards a contradiction that there is an FPTAS $\mathcal{A}$ for f_{SAT}-KCM. We use $\mathcal{A}$ to decide the SAT-matroid problem. Let $I \in \{0,1\}^*$ be a SAT-matroid instance. Consider the following algorithm $\mathcal{B}$ that decides I. For simplicity, let $A = A_I, n = n(A), k = k_I$, and $\alpha = \alpha_I$.

1. Construct the f_{SAT}-KCM instance $U = (I, w_n, w_n, B_{n,k,\alpha})$. Note that $f_{\mathrm{SAT}}(I) = M_{n,k,\alpha}(\mathcal{S}(A))$ by Definition 5.3.
2. Define $\varepsilon = \frac{1}{6 \cdot n^3}$ and compute $\mathcal{A}$ on U and ε; let S be the returned solution.
3. Return that I has a solution if and only if $w_n(S) = B_{n,k,\alpha}$.

By Lemma 5.4, (4), and (5) it follows that U can be constructed from I in time $|I|^{O(1)}$. In addition, $\frac{1}{\varepsilon} = n^{O(1)} = |I|^{O(1)}$; thus, the running time of $\mathcal{B}$ on I is $|I|^{O(1)}$, since $\mathcal{A}$ is an FPTAS. We now show correctness.

- If $\mathcal{B}$ returns that I has a solution. Then, $w_n(S) = B_{n,k,\alpha}$; since S is a solution for U it also holds that $S \in \mathcal{I}_{f_{\mathrm{SAT}}(I)} = \mathcal{I}_{n,k,\alpha}(\mathcal{S}(A))$. Thus, by Corollary 4.4 it holds that I has a solution.
- If I has a solution. Then, there is $T \in \mathcal{L}_{n,k,\alpha}(\mathcal{S}(A))$; thus, by Corollary 4.4 it follows that $w_n(T) = B_{n,k,\alpha}$ implying that T is a solution for U. Therefore,

$$B_{n,k,\alpha} \leq \mathrm{OPT}(U) \leq w_n(S) \leq (1+\varepsilon) \cdot w_n(T) \leq w_n(T) + \frac{1}{6 \cdot n^3} \cdot 3 \cdot n^3 = B_{n,k,\alpha} + \frac{1}{2}. \quad (16)$$

The first inequality holds since the size and the cost of each element are the same, and the size of a solution must be higher or equal from the demand. The third inequality follows because S, T are solutions for U, and S is computed by an FPTAS with the parameter ε. The fourth inequality follows from Lemma 4.1 and by the selection of ε. Observe that the value of S is an integer by (4); thus, by (16), it holds that $w_n(S) = B_{n,k,\alpha}$. Then, $\mathcal{B}$ returns that I has a solution.

Hence, we decide I correctly. It follows that $\mathcal{B}$ is a polynomial algorithm which decides the SAT-matroid problem, contradicting Lemma 5.7. ■

7 Addional Implication

(Ariel :The section is not stable and requires some more work) In this Section we describe some an additional implication of Theorem 1.4 which was omitted from the introduction.

Theorem 1.4 also implies that other results which only apply to representable matroids cannot be generalized to arbitrary matroids. Given a graph $G = (V, E)$ and a path Q in G we use $V(Q)$ to denote the set of vertices on the path Q. Consider for example *generalized weighted colored paths* problem considered by Fomin et al. [10].

Generalized Weighted Colored Paths (GWCP)

Instance	An undirected graph $G = (V, E)$, two sets $A, B \subseteq V$, a representable matroid $(V, \mathcal{I})$ over a finite field of order q (that is, the elements of the matroid are the vertices of the graph), and $p, k, T \in \mathbb{N}$.
A solution	p vertex disjoints paths $Q_1, \ldots, Q_p$ from stating from a vertex in A and terminating in a vertex in B, such that that there is $X \subseteq V(Q_1) \cup \ldots \cup V(Q_p)$ which satisfies $X \in \mathcal{I}$, $w(X) = T$ and $\lvert X \rvert = k$.
Objective	Decide if there is a solution.

In [10] the authors provided an algorithm for GWCP which runs in time $2^{p+O(k^2 \log(q+k))} \cdot n^{O(1)} \cdot T$ where n is the number of vertices of the input graph (Theorem 1.5 in [10]). Consider the following variant of GWCP in which the matroid is arbitrary and given via a membership oracle.

Oracle Generalized Weighted Colored Paths (Oracle-GWCP)

Instance	An undirected graph $G = (V, E)$, two sets $A, B \subseteq V$, a matroid $(V, \mathcal{I})$, and $p, k, T \in \mathbb{N}$.
A solution	p vertex disjoints paths $Q_1, \ldots, Q_p$ from stating from a vertex in A and terminating in a vertex in B, such that that there is $X \subseteq V(Q_1) \cup \ldots \cup V(Q_p)$ which satisfies $X \in \mathcal{I}$, $w(X) = T$ and $\lvert X \rvert = k$.
Objective	Decide if there is a solution.

It can be easily observed that an algorithm for Oracle-GWCP with running time of the form $g(k, p) \cdot (n \cdot T)^{O(1)}$ for an arbitrary function g can be used to decide EMI in time $f(k) \cdot (|I| \cdot T)^{O(1)}$. Given an EMI instance $(E, \mathcal{I}, w, k, T)$ we can generate a graph $G = (E, F)$ where F is a path between two arbitrary elements $a, b \in E$. Also, define $A = \{a\}$ and $B = \{b\}$. It can be easily verified that this Oracle-GWCP has a solution if and only if the EMI instance is a yes instance. Thus, an Oracle-GWCP algorithm in time $g(k, p) \cdot (n \cdot T)^{O(1)}$ implies an EMI algorithm in time $g(k, 1) \cdot (|I| \cdot T)^{O(1)}$, in contradiction to Theorem 1.4.

8 Discussion

References

[1] Berger, A., Bonifaci, V., Grandoni, F., Schäfer, G.: Budgeted matching and budgeted matroid intersection via the gasoline puzzle. Mathematical Programming **128**(1), 355–372 (2011)

[2] Calinescu, G., Chekuri, C., Pal, M., Vondrák, J.: Maximizing a monotone submodular function subject to a matroid constraint. SIAM Journal on Computing **40**(6), 1740–1766 (2011)

[3] Camerini, P.M., Galbiati, G., Maffioli, F.: Random pseudo-polynomial algorithms for exact matroid problems. Journal of Algorithms **13**(2), 258–273 (1992)

[4] Caprara, A., Kellerer, H., Pferschy, U., Pisinger, D.: Approximation algorithms for knapsack problems with cardinality constraints. European Journal of Operational Research **123**(2), 333–345 (2000)

[5] Chakaravarthy, V.T., Choudhury, A.R., Natarajan, S.R., Roy, S.: Knapsack cover subject to a matroid constraint. In: Proc. FSTTCS (2013)

[6] Chekuri, C., Vondrák, J., Zenklusen, R.: Multi-budgeted matchings and matroid intersection via dependent rounding. In: Proc. SODA. pp. 1080–1097 (2011)

[7] Cook, S.A.: The complexity of theorem-proving procedures. In: Logic, Automata, and Computational Complexity: The Works of Stephen A. Cook, pp. 143–152 (2023)

[8] Doron-Arad, I., Kulik, A., Shachnai, H.: An EPTAS for budgeted matching and budgeted matroid intersection via representative sets. In: Proc. ICALP (2023)

[9] Doron-Arad, I., Kulik, A., Shachnai, H.: An EPTAS for budgeted matroid independent set. In: Proc. SOSA. pp. 69–83 (2023)

[10] Fomin, F.V., Golovach, P.A., Korhonen, T., Simonov, K., Stamoulis, G.: Fixed-parameter tractability of maximum colored path and beyond. In: Proc. SODA. pp. 3700–3712 (2023)

[11] Grandoni, F., Zenklusen, R.: Approximation schemes for multi-budgeted independence systems. In: Proc. ESA. pp. 536–548 (2010)

[12] Hassin, R., Levin, A.: An efficient polynomial time approximation scheme for the constrained minimum spanning tree problem using matroid intersection. SIAM Journal on Computing **33**(2), 261–268 (2004)

[13] Lawler, E.L.: Matroid intersection algorithms. Mathematical programming **9**(1), 31–56 (1975)

[14] Lawler, E.L.: Fast approximation algorithms for knapsack problems. Math. Oper. Res. **4**(4), 339–356 (1979)

[15] Lokshtanov, D., Misra, P., Panolan, F., Saurabh, S.: Deterministic truncation of linear matroids. ACM Transactions on Algorithms (TALG) **14**(2), 1–20 (2018)

[16] Oxley, J.G.: Matroid theory, vol. 3. Oxford University Press, USA (2006)

[17] Ravi, R., Goemans, M.X.: The constrained minimum spanning tree problem. In: Algorithm Theory—SWAT'96: 5th Scandinavian Workshop on Algorithm Theory Reykjavík, Iceland, July 3–5, 1996 Proceedings 5. pp. 66–75. Springer (1996)

[18] Schrijver, A.: Combinatorial optimization: polyhedra and efficiency, vol. 24. Springer (2003)